



TiDB as an agentic coding shared memory system

Claude Code, Codex, and Cursor all reading and writing the same TiDB knowledge graph over MCP - so AI-written infrastructure stops duplicating work, breaking conventions, and starting from zero.

Claude Code

Codex

Cursor


Model Context Protocol

Recursive CTE graph

What this is

Teams now write infrastructure with AI coding tools. But each tool works **blind and alone** - it cannot see what already exists, the team's conventions, or how components connect. The result is duplication, drift, broken conventions, and outages.

This demo puts the engineering knowledge where every tool can use it: a **knowledge graph in TiDB Cloud**. Every coding tool connects to it over the **Model Context Protocol (MCP)**, queries it before building, writes its work back, and traverses component relationships with **recursive SQL**. The example workload is an infrastructure-as-code monorepo (Pulumi / TypeScript) for a fictional company, *Acme*, spanning AWS, Cloudflare, and Okta.



Three real CLIs → two MCP servers → one TiDB database. The dashboard reads the same database live.

Proven on real TiDB - every screenshot below is from an actual run. The graph, the writes from three real CLIs (Claude Code, Codex, Cursor), and the recursive-CTE traversals all ran against a live TiDB Cloud cluster (v8.5.3) over the MySQL protocol with TLS. The terminal captures are the unedited tool output.

The problem: AI tools with no shared memory

Four failure modes show up the moment more than one agent (or session) touches the same codebase.

FAILURE MODE	WHAT HAPPENS WITHOUT A SHARED GRAPH
Duplication	A tool can't see a component already exists, so it builds a second one. Drift and double cost follow.
Convention breaks	It writes a raw <code>aws.s3.Bucket</code> instead of the approved library wrapper, skipping required tags and policies.
Invisible blast radius	It changes a shared library and breaks resources two hops away - discovered only in production.
Cold starts	The next tool or session re-derives everything from the filesystem and re-makes the same mistakes.

The common root cause: the knowledge that prevents these - **what exists, how it connects, and the rules** - lives only in scattered code and people's heads. It is not queryable, and it is not shared across tools.

The fix. Model the codebase as a graph in TiDB: *infra_components* (nodes), *component_edges* (typed relationships), and *session_log* (who did what). Expose it to every tool over MCP. Now the tools query before they build and hand off warm.

The live demo: three tools, one task

The task: production is complete, but **staging is missing its DNS and SSO layer** (four components). Watch three different AI coding tools finish it - each picking up where the last left off, because they share one TiDB brain. The terminal captures below are unedited tool output.

Step 1 - Claude Code inspects, then builds the asset layer claude-code

It first **queries** the graph (it does not assume), then scaffolds the missing static-assets bucket and its Cloudflare DNS record by composing the existing *S3Bucket* and *DnsRecord* libraries - never raw cloud resources - and **writes each one back**.

```
• Both components recorded (IDs 25 and 26). Here's a summary of everything done:

---
Gap found - staging missing vs production:



| Resource                             | Status       |
|--------------------------------------|--------------|
| acme-staging-static-assets (S3)      | Scaffolded   |
| acme-staging-assets-dns (Cloudflare) | Scaffolded   |
| acme-staging-admin-dns (Cloudflare)  | Out of scope |
| acme-staging-admin-sso (Okta)        | Out of scope |



Files created:
- environments/staging/storage.ts - adds acme-staging-static-assets via S3Bucket (existing acme-staging-data-exports retained)
- environments/staging/dns.ts - adds acme-staging-assets-dns via DnsRecord, wired to staticAssets.bucketDomainName

Rules followed:
- No raw aws.* resources - both use the library wrappers (S3Bucket, DnsRecord)
- proxied: true is library-enforced on DnsRecord (never overridden)
- All four required tags present: Environment, ManagedBy, Component, Team
- Physical names follow acme-<env>-<name> convention
- Both components recorded in the KB with depends_on edges to their respective libraries

* Baked for 2m 14s
```

Real Claude Code output. It recorded both components (IDs 25 & 26) and reports the rules it followed: no raw *aws.**, *proxied: true* library-enforced, all four required tags, *acme-<env>-<name>* naming, KB writeback with dependency edges.

Step 2 - Codex picks up warm codex

Codex opens with **zero local memory** of Claude Code's work. It reads the shared *session_log*, sees what was just created, and continues with the admin DNS record. No copy-pasting context between tools.

```
• I have enough context to continue the build consistently. I'm recording acme-staging-admin-dns now as a staging Cloudflare component that composes DnsRecord, with the summary mirrored from production and adjusted for staging.

• Called
  tidb-infra-kb.write_component({"depends_on":"DnsRecord","developer":"codex","env":"staging","name":"acme-staging-admin-dns","summary":"Cloudflare DNS record for the admin portal. acme-staging-admin-sso redirects here.","type":"Cloudflare"})
  Written: acme-staging-admin-dns (id=27)

• The previous session most recently created acme-staging-assets-dns, immediately after acme-staging-static-assets.

I continued the staging build by recording acme-staging-admin-dns as a Cloudflare component in staging, depending on DnsRecord, with the production-matching summary adapted for staging. write_component succeeded with id 27.

- Worked for 5m 18s -
```

Real Codex output: "The previous session most recently created *acme-staging-assets-dns*" - it read the shared log, then wrote *acme-staging-admin-dns* (id=27). A different vendor's tool, continuing seamlessly.

Step 3 - Cursor verifies with a recursive query, then finishes cursor

Before writing the SSO application, Cursor runs a **recursive query** to confirm the dependency pattern (an SSO app must redirect to a DNS record), then scaffolds it correctly. Staging reaches parity.

```
Recursive CTE: acme-prod-admin-sso transitive dependencies

Ran the project's downstream CTE via db_query:

| Depth | From | Relationship | To | Type |
| 1 | acme-prod-admin-sso | uses | SsoApplication | Library |
| 1 | acme-prod-admin-sso | redirects_to | acme-prod-admin-dns | Cloudflare |
| 2 | acme-prod-admin-dns | uses | DnsRecord | Library |

Transitive closure: SsoApplication → acme-prod-admin-dns → DnsRecord.

SSO must redirect to a DnsRecord

Confirmed in the graph and conventions:

• Production has acme-prod-admin-sso redirects_to acme-prod-admin-dns, and that DNS record uses DnsRecord.
• SsoApplication is documented as requiring redirect URIs on the matching DnsRecord hostname (libs/sso-app, configs/claude-code/CLAUDE.md).

So the SSO app does not redirect to a raw hostname in isolation – it redirects to the admin DNS component, which composes DnsRecord.


Staging scaffold + KB writeback

1. environments/staging/dns.ts – added acme-staging-admin-dns and adminDnsHostname (needed for the redirect URI).
2. environments/staging/sso.ts – new SsoApplication with redirect https://{adminDnsHostname}/callback.
3. write_component – recorded acme-staging-admin-sso (Okta, staging, composes SsoApplication).
4. db_execute – edge acme-staging-admin-sso redirects_to acme-staging-admin-dns.

KB edges for staging SSO:

• uses → SsoApplication
• redirects_to → acme-staging-admin-dns
```


Real Cursor output - the headline moment. The recursive query returns the chain SsoApplication → acme-prod-admin-dns → DnsRecord, proving the dependency before writing a line of code.



The result, in the dashboard: all three tools' activity in one feed (cursor, codex, claude-code) and the staging-parity strip now fully green. Built by three tools sharing one memory.

The query a vector database can't answer

The headline capability: **graph reachability**. "What breaks if I change this component?" requires traversing dependency edges to arbitrary depth. A recursive CTE does it in one query.




Blast radius of the S3Bucket library on the live dashboard, visualised as a depth-coloured subgraph with the exact SQL and result rows beside it. The traversal reaches the RDS databases **two hops away** - their backup buckets compose S3Bucket - a connection no text search could surface.

```
-- What breaks if S3Bucket changes? Runs unchanged on SQLite and TiDB.
WITH RECURSIVE blast(from_id, to_id, relationship, depth) AS (
  SELECT e.from_id, e.to_id, e.relationship, 1
  FROM component_edges e JOIN infra_components c ON e.to_id = c.id
  WHERE c.name = 'S3Bucket'
  UNION ALL
  SELECT e.from_id, e.to_id, e.relationship, b.depth + 1
  FROM component_edges e JOIN blast b ON e.to_id = b.from_id
  WHERE b.depth < 10
)
SELECT depth, from_id, relationship, to_id FROM blast ORDER BY depth;
```

Why a vector store can't do this. Vector similarity finds *related-looking* text. It cannot compute a transitive closure over a dependency graph. This is reachability - exactly what SQL recursion is built for - and TiDB runs it as a distributed query alongside the same tables that hold the components, their tags, and (optionally) their vector embeddings.


One engine: vector + full-text + graph, one table

The same *infra_components* table that holds the graph also carries a 1024-dim *VECTOR* column - embedded server-side by TiDB's *EMBED_TEXT* - and a *FULLTEXT* index. Semantic, keyword, and hybrid search all run on it. No separate vector database, no ETL, no sync lag.




Hybrid search for "admin portal access control login". Reciprocal rank fusion over two queries on the same table. The top two rows match on **both** signals (vector + FTS ✓); the rest are **vector-only** - the keyword index missed them (FTS ✗) but the embedding caught them semantically.

The headline: full-text alone returns **2** results; the same query, same table, via embeddings returns **6** - including *SsoApplication* and the analytics database, which contain none of the query's keywords. Hybrid gives you both, ranked together.



Vector mode: pure semantic ranking by cosine distance (*VEC_COSINE_DISTANCE* over *EMBED_TEXT*). 6 results.



Full-text mode: keyword ranking via *FTS_MATCH_WORD* on the *FULLTEXT* index. Only 2 results - keywords miss the rest.

Why it matters: a coding agent searching its memory for "how do we do auth" needs *meaning*, not string matching. TiDB gives the agent vector recall, keyword precision, and graph reachability through one SQL endpoint - the same endpoint that already holds the structured state.

Why TiDB is the right brain for agents

In the emerging pattern, the database is the agent's persistent memory, execution-state store, and shared knowledge base. That demands more than a vector index.

One engine, every access pattern

Agents need **SQL** for structured state, **graph traversal** for relationships, **vector search** for semantic memory, and **full-text (BM25)** for knowledge - all over the same data, with no ETL and no sync lag. TiDB provides all of them in one MySQL-compatible engine. This demo uses SQL + recursive CTEs today; the same *infra_components* table can carry a *VECTOR* column for hybrid graph-plus-semantic retrieval with *VEC_COSINE_DISTANCE*.

ACID for agent state

Each write here is one transaction: **component + dependency edge + audit-log entry, all or nothing**. TiDB's distributed ACID prevents the partial-write bugs that make multi-agent systems nearly impossible to debug.

Built for agent scale

TiDB Cloud scales to zero with pay-per-request pricing. Millisecond database branching gives each agent or workspace an isolated environment - the pattern behind production fleets running millions of agent databases on TiDB Cloud.

How it compares for this workload

CAPABILITY	TIDB CLOUD	VECTOR DB (PINECONE)	POSTGRES / PGVECTOR
Recursive graph traversal (CTE)	✓	✗	✓ (single node)
SQL + JOINS + ACID	✓	✗	✓
Vector + full-text in one engine	✓	vector only	partial
Horizontal scale / millions of tenants	✓	partial	✗

Proof at scale: a leading agent platform runs 10M+ databases and billions of events/day on TiDB Cloud, with 100% of those databases created by agents.

Concrete scenarios the graph prevents

Two everyday infrastructure mistakes, and how a queryable graph stops them before code is written.

Infra Knowledge Base

SQL Cloud Services

Clouds Data

Code

Control

+ new TiDB

Live

Dependency graph (CTE)

Why TiDB - scenarios

Reset KB

The duplicate backup bucket

cursor

HEADLINE

Task: Add backups for the staging analytics database.

WITHOUT THE GRAPH

WITHOUT TiDB KNOWLEDGE GRAPH

THE QUERY THAT ANSWERS IT

Why TiDB: The composition relationship lives in the graph, not buried in Typescript. A flat repo search or a vector store of code snippets can't tell you "this library already creates that resource for you".

Blast radius before a refactor

HEADLINE

Task: Change the naming scheme inside the S3Bucket library.

WITHOUT THE GRAPH

WITHOUT TiDB KNOWLEDGE GRAPH

THE QUERY THAT ANSWERS IT

Why TiDB: A recursive CTE traverses arbitrary-depth dependency chains in one query. Vector similarity finds "related-looking" code; it can't compute a transitive closure. This is graph reachability - exactly what TiDB recursion is for.

Staging parity, picked up warm across tools

HEADLINE

Task: Bring staging to parity with production DNS + SSO.

WITHOUT THE GRAPH

WITHOUT TiDB KNOWLEDGE GRAPH

THE QUERY THAT ANSWERS IT

Why TiDB: The session log = component graph are shared data every tool reads and writes. The next session starts warm instead of re-deriving context from the filesystem.

In-product scenario cards. Each contrasts the failure without the graph against the corrected path with TiDB, and shows the exact query that answers it.

The duplicate-bucket guardrail, caught live

Asked to add a backup bucket for the staging database, a coding tool queried the graph first - and stopped itself:

```

> I'm about to add a backup S3 bucket for acme-staging-analytics-db. First query the KB: does PostgresDatabase already create one?
• No need to re-query - the KB results from this session are definitive:

Yes, PostgresDatabase already creates one automatically.

From the library summary:
| "Automatically instantiates an S3Bucket for backups - do NOT create a separate backup bucket manually."

The edge confirms it: PostgresDatabase --instantiates--> S3Bucket.

And the production precedent shows the naming: acme-prod-analytics-db auto-created acme-prod-analytics-db-backup. Staging would
follow the same pattern, so acme-staging-analytics-db-backup already exists as a side-effect of the acme-staging-analytics-db
instantiation.

Don't add a manual backup bucket - you'd end up with a duplicate, and the library's lifecycle rules/tags on the auto-created one
would conflict with yours.

```

Real tool output: it found PostgresDatabase --instantiates—> S3Bucket in the graph and concluded "Don't add a manual backup bucket - you'd end up with a duplicate, and the library's lifecycle rules/tags on the auto-created one would conflict with yours." The duplicate was never written.

The duplicate backup bucket. Asked to "add backups" for a database, a tool greps, finds nothing, and creates a raw duplicate. The graph shows *PostgresDatabase* already *instantiates* a backup bucket - so the right action is to write nothing.

Blast radius before a refactor. Before changing the *S3Bucket* library, one recursive query returns the full set of dependents across every environment - including the RDS instances two hops away. The change is staged safely instead of breaking production.

How the tools are driven: the prompts

The behaviour is prompt-driven. Two layers: a **standing instruction** in each tool's rules file (always read/write the shared memory), and the **per-task prompt** the developer types. Nothing is hard-coded - this is how a real team would configure it.

Standing instruction (rules file: CLAUDE.md / AGENTS.md / .cursorrules)

```
Before writing any infrastructure code, query the tidb-infra-kb MCP first
(what exists, how it connects, what previous sessions did). Never use raw
aws.* / cloudflare.* resources - compose the library components. After you
create anything, record it with write_component so the next session starts warm.
```

Pane 1 - Claude Code

```
Use the tidb-infra-kb MCP. Query the knowledge base: what components exist in
staging vs production, and what is staging missing? Then scaffold the missing
staging static-assets bucket and its Cloudflare DNS record, composing the
S3Bucket and DnsRecord libraries - never raw aws.* resources. Follow the
acme-<env>-<name> naming and required tags. Record each new component with
write_component.
```

Pane 2 - Codex (opens cold, picks up warm)

```
Use the tidb-infra-kb MCP. Read the recent session_log - what did the previous
session just create? Continue the staging build: add the admin-portal DNS record
(acme-staging-admin-dns) matching the production pattern, composing DnsRecord.
Record it with write_component.
```

Pane 3 - Cursor (verify with a recursive CTE, then build)

```
Use the tidb MCP (db_query) to run a recursive CTE: what does the production
acme-prod-admin-sso transitively depend on? Confirm an SSO app must redirect to a
DnsRecord. Then scaffold acme-staging-admin-sso composing SsoApplication, with the
redirect URI pointing at acme-staging-admin-dns. Record it with write_component.
```

The guardrail prompt (duplicate-bucket scenario)

```
I'm about to add a backup S3 bucket for acme-staging-analytics-db. First query
the KB: does PostgresDatabase already create one?
```

Every query the demo runs

Everything below executes on TiDB Cloud over the MySQL protocol - either through the official *tidb* MCP (*db_query* / *db_execute*) or the convention *tidb-infra-kb* MCP. Nothing is mocked.

1. Inspect - what exists where (Claude Code)

```
SELECT environment, name, component_type
FROM infra_components
WHERE environment IN ('staging', 'production')
ORDER BY environment, component_type;
```

2. Pick up warm - read the cross-session log (Codex)

```
SELECT developer, action, detail, created_at
FROM session_log
ORDER BY created_at DESC LIMIT 5;
```

3. Prove a dependency pattern - recursive CTE (Cursor)

```
-- what does acme-prod-admin-sso transitively depend on?
WITH RECURSIVE deps(from_id, to_id, relationship, depth) AS (
  SELECT e.from_id, e.to_id, e.relationship, 1
  FROM component_edges e JOIN infra_components c ON e.from_id = c.id
  WHERE c.name = 'acme-prod-admin-sso'
  UNION ALL
  SELECT e.from_id, e.to_id, e.relationship, d.depth + 1
  FROM component_edges e JOIN deps d ON e.from_id = d.to_id
  WHERE d.depth < 10
)
SELECT d.depth, cf.name, d.relationship, ct.name
FROM deps d JOIN infra_components cf ON d.from_id = cf.id
      JOIN infra_components ct ON d.to_id = ct.id
ORDER BY d.depth;
```

4. The atomic write - what write_component actually does

```
-- one transaction: component + dependency edge + audit log, all or nothing.
-- the embedding is generated server-side by TiDB, no client model.
INSERT INTO infra_components (name, component_type, environment, summary, embedding)
VALUES ('acme-staging-admin-sso', 'Okta', 'staging', '...',
      EMBED_TEXT('tidbcloud_free/amazon/titan-embed-text-v2', '...'));
INSERT INTO component_edges (from_id, to_id, relationship)
VALUES (@sso_id, @dns_id, 'redirects_to');
INSERT INTO session_log (developer, action, detail)
VALUES ('cursor', 'created', 'Created Okta component acme-staging-admin-sso');
```

5. Blast radius - recursive CTE (Dependency-graph tab)

```
-- what breaks if S3Bucket changes? walks dependents upward, any depth.
WITH RECURSIVE blast(from_id, to_id, relationship, depth) AS (
  SELECT e.from_id, e.to_id, e.relationship, 1
```

```

FROM component_edges e JOIN infra_components c ON e.to_id = c.id
WHERE c.name = 'S3Bucket'
UNION ALL
SELECT e.from_id, e.to_id, e.relationship, b.depth + 1
FROM component_edges e JOIN blast b ON e.to_id = b.from_id
WHERE b.depth < 10
)
SELECT DISTINCT depth, from_id, relationship, to_id FROM blast ORDER BY depth;

```

6. Search - vector, full-text, and hybrid (Search tab)

```

-- semantic: TiDB embeds the query, ranks by cosine distance
SELECT name, VEC_COSINE_DISTANCE(embedding,
    EMBED_TEXT('tidbcloud_free/amazon/titan-embed-text-v2', 'admin portal access control'))
FROM infra_components ORDER BY distance LIMIT 6;

-- keyword: BM25 over the FULLTEXT index
SELECT name FROM infra_components
WHERE FTS_MATCH_WORD('admin portal access control', summary)
ORDER BY FTS_MATCH_WORD('admin portal access control', summary) DESC LIMIT 6;

-- hybrid = reciprocal rank fusion of the two result sets, same table.

```

Running it on your own repo

Everything you saw - query-before-build, convention checks, recursive dependency traversal, vector + full-text search, cross-tool handoff - is already general. The one piece to add for a real codebase is a small **ingestion job** that populates the graph from your source. Everything downstream is unchanged; it reads the same three tables.

The ingestion job (runs in CI on every merge)

- 1 **Read the resource graph.** Parse the Pulumi program (`pulumi stack export` / the resource tree, or the TypeScript AST) to enumerate every component and its parent-child and cross-references.
- 2 **Upsert nodes + edges.** Each component becomes a row in `infra_components`; each relationship a row in `component_edges` (`instantiates`, `fronts`, `redirects_to`, ...).
- 3 **Embed.** Call `EMBED_TEXT` on each summary so semantic search works the moment it lands - no embedding pipeline to run.
- 4 **Log the scan** in `session_log` so tools can see "graph last synced at commit abc123."

```
# ingest.py - run in CI; idempotent upserts keep the graph in sync with main
for resource in pulumi_resource_graph():
    write_component(name=resource.urn_name, type=resource.kind, env=resource.stack,
                   summary=resource.doc, depends_on=resource.parent) # embeds + edges, atomic
```

That is the whole gap between this demo and production. The graph the tools query is the same shape; only its *source* changes from the seed script to your live repo. Query, CTE, search, convention-enforcement, and cross-tool memory all work as shown.

How to run the demo

The repository ships everything: the TiDB-backed knowledge base, both MCP servers, per-tool configs, the dashboard, and a presenter's runbook.

- 1 **Configure TiDB.** `cp .env.example .env` and paste your TiDB Cloud credentials (Connect → General).
- 2 **Set up.** `./setup.sh` - installs dependencies, creates the schema and seeds the graph on TiDB, and generates MCP configs for Claude Code, Codex, and Cursor.
- 3 **Launch the dashboard.** `./demo.sh` → open `http://localhost:7001` (the screen you narrate).
- 4 **Open three terminals.** In each, run `claude`, `codex`, and `cursor-agent` from the repo. All three load both MCP servers automatically.
- 5 **Follow the prompts** in `DEMO.md` / `PRESENTER.md`. As the tools work, the dashboard updates live from TiDB.

The two MCP servers each tool connects to

SERVER	ORIGIN	WHAT IT PROVIDES
<code>tidb</code>	PingCAP official	<code>db_query</code> / <code>db_execute</code> - raw SQL, recursive CTEs, vector search
<code>tidb-infra-kb</code>	This repo	<code>query_knowledge_base</code> + <code>write_component</code> - convention-checked, atomic writes

Repository: github.com/stephenlthorn/mem9-ai-coding · Apache-2.0. The dashboard falls back to local SQLite if no TiDB credentials are present, so the demo runs even with no network.